Department of Statistics, The Chinese University of Hong Kong
STAT5010 Advanced Statistical Inference | Term 1, 2019–20

Take-home Examination

Instruction to the candidates: Please attempt all of the questions. Each problem carries an equal weight of 4 points. Your final score will be capped by 20, which is also the defined full mark of this exam. Good luck!

1. Let $X_1, \ldots X_n$ be a random sample from a $N(\theta, \sigma^2)$ population with σ^2 known. Consider estimating θ using the squared error loss. Let $\pi(\theta)$ be a $N(\mu, \tau^2)$ prior distribution on θ and let δ^π be the Bayes estimator of θ. Verify the following formulas for the risk function and Bayes risk.

   (a) For any constants a and b, the estimator $\delta(\boldsymbol{X}) = a\bar{\boldsymbol{X}} + b$ has risk function
   $$R(\theta, \delta) = a^2 \frac{\sigma^2}{n} + \{b - (1-a)\theta\}^2.$$

   (b) Let $\eta = \sigma^2/(n\tau^2 + \sigma^2)$. The risk function for the Bayes estimator is
   $$R(\theta, \delta^\pi) = (1-\eta)^2 \frac{\sigma^2}{n} + \eta^2(\theta - \mu)^2.$$

   (c) The Bayes risk for the Bayes estimator is
   $$B(\pi, \delta^\pi) = \tau^2 \eta.$$

2. Let X be an observation from the pdf
   $$f(x \mid \theta) = \left(\frac{\theta}{2}\right)^{|x|} (1-\theta)^{1-|x|}, \quad x \in \{-1, 0, 1\}; \theta \in [0, 1].$$

   (a) Find the MLE of θ.

   (b) Define an estimator $T(X)$ by
   $$T(X) = \begin{cases} 2 & \text{, if } x = 1 \\ 0 & \text{, otherwise} \end{cases}.$$
   Show that $T(X)$ is an unbiased estimator of θ.

   (c) Find a better estimator than $T(X)$ and prove that it is better.

3. Consider a Bayesian model in which the prior distribution for Θ is standard exponential and the density for X given Θ is
   $$f(x \mid \theta) = e^{\theta - x} I(x > \theta).$$

   (a) Find the marginal density for X and $E(X)$ in the Bayesian model.

   (b) Find the Bayes estimator for Θ under squared error loss. (Assume $X > 0$.)

4. Let $X_1, X_2, \ldots, X_n$ be i.i.d. from the uniform distribution on $(1, 2)$, and let H_n denote the harmonic average of the first n variables:

$$H_n = \frac{n}{X_1^{-1} + \ldots + X_n^{-1}}.$$

   (a) Show that $H_n \xrightarrow{p} c$ as $n \to \infty$, identifying the constant c.

   (b) Show that $\sqrt{n}(H_n - c)$ converges in distribution, and identify the limit.

5. Let $X_1, \ldots, X_n$ be i.i.d. from $N(\theta, 1)$ and let $U_1, \ldots, U_n$ be i.i.d. from a uniform distribution on $(0, 1)$, with all $2n$ variables independent. Define $Y_i = X_i U_i$, $i = 1, \ldots, n$. If the X_i and U_i are both observed, then $\bar{X}$ would be a natural estimator for θ. If only the products $Y_1, \ldots, Y_n$ are observed, then $2\bar{Y}$ may be a more responsible estimator. Determine the asymptotic relative efficiency (ARE) of $2\bar{Y}$ with respect to $\bar{X}$, where ARE of $\hat{\theta}_n$ with respect to $\tilde{\theta}_n$ is defined as the ratio $\sigma^2_{\tilde{\theta}}/\sigma^2_{\hat{\theta}}$ if $\sqrt{n}(\hat{\theta} - \theta_0) \xrightarrow{d} N(0, \sigma^2_{\hat{\theta}})$ and $\sqrt{n}(\tilde{\theta} - \theta_0) \xrightarrow{d} N(0, \sigma^2_{\tilde{\theta}})$, respectively.

6. Suppose $X_1, \ldots, X_n \overset{i.i.d.}{\sim} N(\theta, 1)$ for some $\theta \in \mathbb{R}$, and we want to estimate θ with respect to squared error loss. Show that the estimator $\delta_a(X) = \bar{X} + a$ is not a Bayes estimator for any a.

- End -